

Project: Exploring Food Nutrition Patterns

Objective: Analyze the nutritional content of various foods using the CORGIS Food dataset. The purpose is to discover patterns and answer key questions related to protein, fat, and carbohydrate content across different foods and categories, using data science tools and techniques.

Analytical Questions:

1. Which foods have the highest and lowest protein, fat, and carbohydrate content?
2. How do average nutrient values differ among the most common food categories?
3. What relationships exist between foods' protein and fat content?
4. Can similar foods be grouped by their nutrient profiles using clustering?

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
data = pd.read_csv('food.csv')
data.head()
```

Out[1]:

	Category	Description	Nutrient				
			Data Bank Number	Data.Alpha Carotene	Data.Beta Carotene	Data.Beta Cryptoxanthin	Data.Carbohydr
0	Milk	Milk, human	11000000	0	7	0	€
1	Milk	Milk, NFS	11100000	0	4	0	₺
2	Milk	Milk, whole	11111000	0	7	0	₺
3	Milk	Milk, low sodium, whole	11111100	0	7	0	₺
4	Milk	Milk, calcium fortified, whole	11111150	0	7	0	₺

5 rows × 38 columns



Data Acquisition and Preparation

I used the CORGIS Food dataset, which provides nutrient information for thousands of foods. After loading the data, I performed basic cleaning, selecting relevant columns, checking for

missing values, and focusing on the ten most common food categories to ensure clear visualization and comparison.

```
In [2]: print(data.columns)
```

```
Index(['Category', 'Description', 'Nutrient Data Bank Number',
       'Data.Alpha Carotene', 'Data.Beta Carotene', 'Data.Beta Cryptoxanthin',
       'Data.Carbohydrate', 'Data.Cholesterol', 'Data.Choline', 'Data.Fiber',
       'Data.Lutein and Zeaxanthin', 'Data.Lycopene', 'Data.Niacin',
       'Data.Protein', 'Data.Retinol', 'Data.Riboflavin', 'Data.Selenium',
       'Data.Sugar Total', 'Data.Thiamin', 'Data.Water',
       'Data.Fat.Monosaturated Fat', 'Data.Fat.Polysaturated Fat',
       'Data.Fat.Saturated Fat', 'Data.Fat.Total Lipid',
       'Data.Major Minerals.Calcium', 'Data.Major Minerals.Copper',
       'Data.Major Minerals.Iron', 'Data.Major Minerals.Magnesium',
       'Data.Major Minerals.Phosphorus', 'Data.Major Minerals.Potassium',
       'Data.Major Minerals.Sodium', 'Data.Major Minerals.Zinc',
       'Data.Vitamins.Vitamin A - RAE', 'Data.Vitamins.Vitamin B12',
       'Data.Vitamins.Vitamin B6', 'Data.Vitamins.Vitamin C',
       'Data.Vitamins.Vitamin E', 'Data.Vitamins.Vitamin K'],
      dtype='object')
```

Exploratory Analysis: Nutrient Extremes

I identified the foods with the highest and lowest values for protein, fat, and carbohydrate. This comparison reveals which foods are most and least nutrient-dense and highlights differences in macronutrient profiles across products.

```
In [3]: # Top 10 highest protein foods
protein_sorted = data[['Description', 'Data.Protein']].sort_values(by='Data.Protein')
print('Top 10 highest protein foods:')
print(protein_sorted.head(10))

print('Bottom 10 lowest protein foods:')
print(protein_sorted.tail(10))
```

Top 10 highest protein foods:

	Description	Data.Protein
7022	Nutritional powder mix, whey based, NFS	78.13
7025	Nutritional powder mix, protein, NFS	78.13
1528	Tuna, fresh, dried	76.25
7010	Nutritional powder mix (EAS Whey Protein Powder)	66.67
1468	Salmon, dried	64.06
1306	Cod, dried, salted	62.82
1249	Fish, NS as to type, dried	62.82
826	Pork skin rinds	61.30
1584	Squid, dried	58.94
7013	Nutritional powder mix (Isopure)	58.14

Bottom 10 lowest protein foods:

	Description	Data.Protein
6880	Fruit flavored drink, powdered, not reconstituted	0.0
6882	Sports drink, dry concentrate, not reconstituted	0.0
6351	Olive oil	0.0
6742	Coffee, instant, decaffeinated, pre-lightened ...	0.0
6352	Peanut oil	0.0
6353	Canola oil	0.0
6348	Corn oil	0.0
6888	Alcoholic malt beverage, sweetened	0.0
6889	Alcoholic malt beverage, higher alcohol, sweet...	0.0
7082	Industrial oil as ingredient in food	0.0

```
In [4]: print(data['Category'].unique()[:20]) # See the first 20 unique values
print(len(data['Category'].unique())) # Total number of unique categories
```

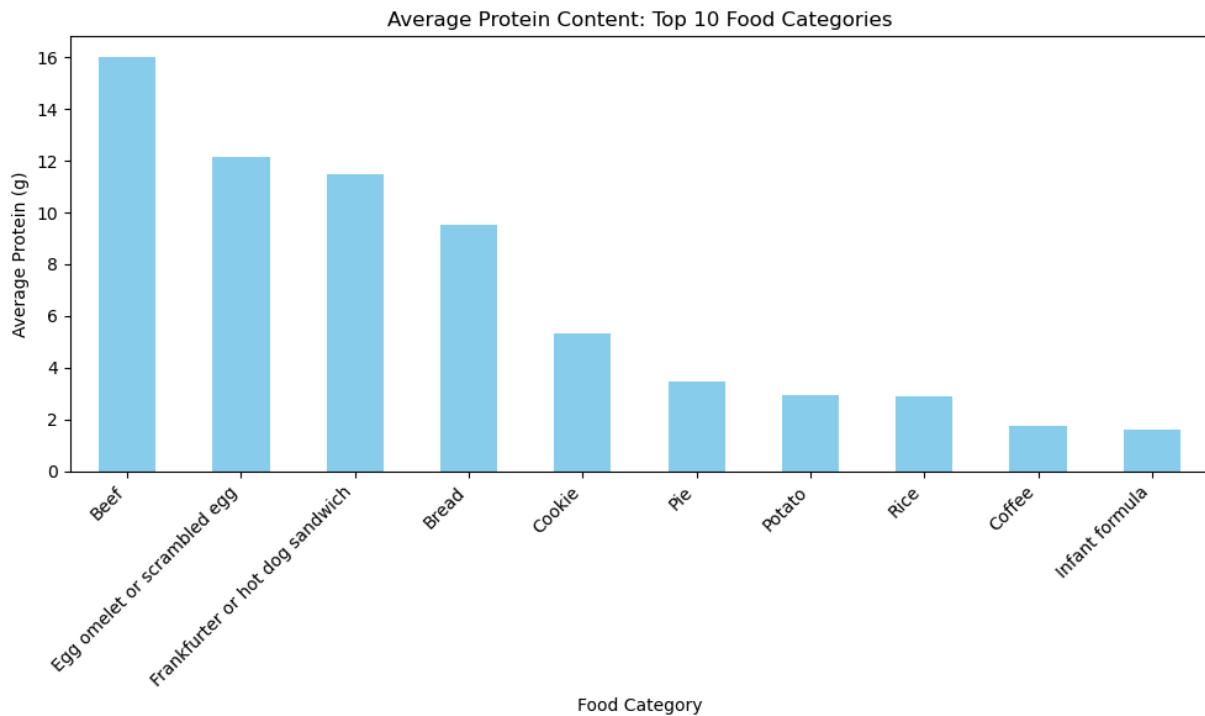
```
['Milk' 'Buttermilk' 'Kefir' "Goat's milk" 'Non-dairy milk' 'Soy milk'
 'Almond milk' 'Rice milk' 'Coconut milk' 'Yogurt' 'Chipotle dip'
 'Dill dip' 'Onion dip' 'Ranch dip' 'Spinach dip' 'Tzatziki dip'
 'Vegetable dip' 'Yogurt parfait' 'Frozen yogurt' 'Frozen yogurt sandwich']
```

2429

```
In [5]: # Find the top 10 most common categories
top_categories = data['Category'].value_counts().head(10).index
subset = data[data['Category'].isin(top_categories)]

# Calculate average protein for just these groups
avg_protein = subset.groupby('Category')['Data.Protein'].mean().sort_values(ascending=False)

plt.figure(figsize=(10,6))
avg_protein.plot(kind='bar', color='skyblue')
plt.title('Average Protein Content: Top 10 Food Categories')
plt.ylabel('Average Protein (g)')
plt.xlabel('Food Category')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.savefig('test.png')
plt.show()
```



Interpretation: Average Protein by Food Category

The bar chart above shows the average protein content for the 10 most frequent food categories in the dataset. It is clear that animal-based protein sources like beef and eggs dominate the top ranks for protein. Processed foods like frankfurters and bread have moderate protein, while baked goods, starchy foods, and beverages have much lower averages. This visualization makes it easy to compare broad nutritional patterns across everyday foods.

```
In [6]: # Sort foods by total fat content
totfat_sorted = data[['Description', 'Data.Fat.Total Lipid']].sort_values(by='Data.
print('Top 10 highest fat foods:')
print(totfat_sorted.head(10))
```

```
print('Bottom 10 lowest fat foods:')
print(totfat_sorted.tail(10))
```

Top 10 highest fat foods:

		Description	Data.Fat.Total	Lipid
7082	Industrial oil as ingredient in food		100.0	
6354	Safflower oil		100.0	
6345	Vegetable oil, NFS		100.0	
6359	Wheat germ oil		100.0	
6358	Walnut oil		100.0	
6357	Sunflower oil		100.0	
6356	Soybean oil		100.0	
6346	Almond oil		100.0	
6355	Sesame oil		100.0	
6353	Canola oil		100.0	

Bottom 10 lowest fat foods:

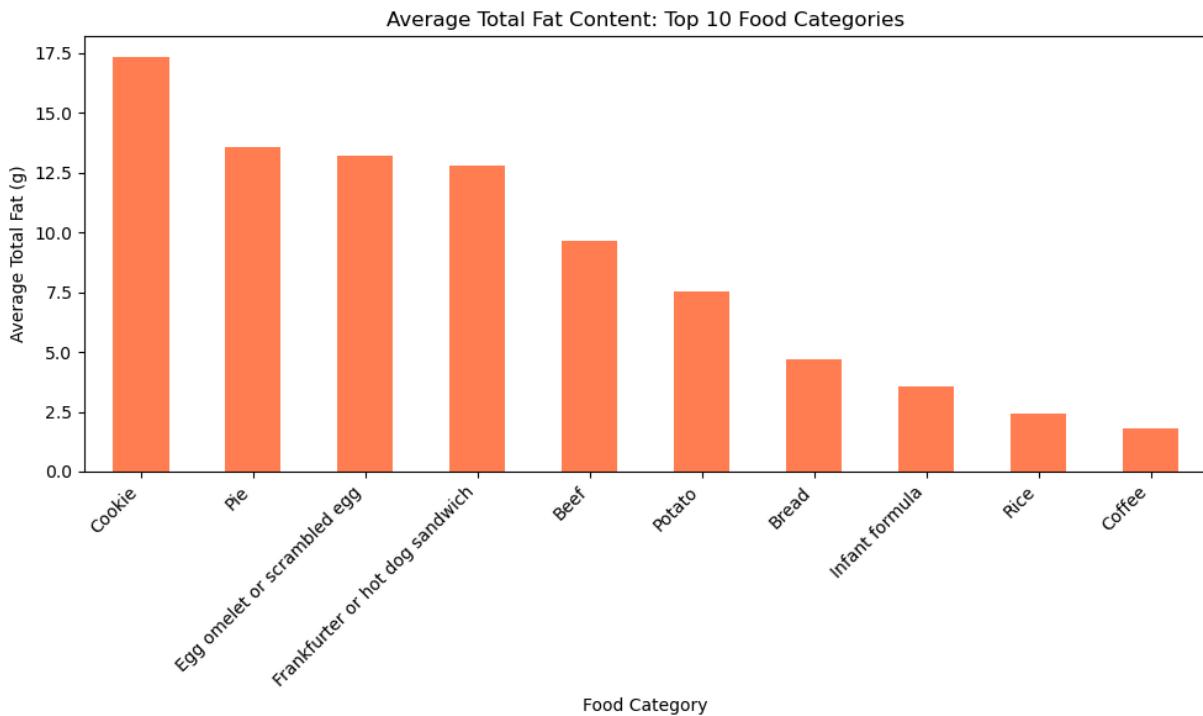
		Description	Data.Fat.Total	Lipid
6859	Vegetable and fruit juice drink, with high vit...		0.0	
104	Chocolate milk, ready to drink, fat free (Nesq...		0.0	
6848	Fruit flavored drink, with high vitamin C, pow...		0.0	
6856	Lemonade, fruit juice drink, light		0.0	
6440	Blueberry syrup		0.0	
6854	Orange juice beverage, 40-50% juice, light		0.0	
6852	Cranberry juice drink, with high vitamin C, light		0.0	
6851	Fruit juice drink, diet		0.0	
6849	Fruit juice drink, with high vitamin C, light		0.0	
6819	Soft drink, chocolate flavored		0.0	

```
In [7]: top_categories = data['Category'].value_counts().head(10).index
subset = data[data['Category'].isin(top_categories)]

avg_fat = subset.groupby('Category')['Data.Fat.Total Lipid'].mean().sort_values(ascending=False)

plt.figure(figsize=(10,6))
avg_fat.plot(kind='bar', color='coral')
plt.title('Average Total Fat Content: Top 10 Food Categories')
plt.ylabel('Average Total Fat (g)')
plt.xlabel('Food Category')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()

plt.show()
```



```
In [8]: # Sort foods by carbohydrate
totcarb_sorted = data[['Description', 'Data.Carbohydrate']].sort_values(by='Data.Carbohydrate')
print('Top 10 highest carbohydrate foods:')
print(totcarb_sorted.head(10))

print('Bottom 10 lowest carbohydrate foods:')
print(totcarb_sorted.tail(10))
```

Top 10 highest carbohydrate foods:

	Description	Data.Carbohydrate
6429	Sugar substitute, stevia, powder	100.00
6431	Sugar substitute, monk fruit, powder	100.00
6424	Sugar, white, confectioner's, powdered	99.77
6423	Sugar, white, granulated or lump	99.60
6422	Sugar, NFS	99.60
6427	Sugar substitute and sugar blend	99.35
372	Strawberry beverage powder, dry mix, not recon...	99.10
6879	Fruit flavored drink, with high vitamin C, pow...	98.94
6602	Gumdrops	98.90
6617	Dietetic or low calorie candy, NFS	98.60

Bottom 10 lowest carbohydrate foods:

	Description	Data.Carbohydrate
658	Beef steak, NS as to cooking method, lean and ...	0.0
976	Chicken leg, drumstick and thigh, fried, coate...	0.0
657	Beef steak, NS as to cooking method, NS as to ...	0.0
656	Beef, NS as to cut, fried, NS to fat eaten	0.0
6625	Coffee, NS as to type	0.0
6626	Coffee, NS as to brewed or instant	0.0
6627	Coffee, brewed	0.0
6628	Coffee, brewed, blend of regular and decaffein...	0.0
974	Chicken leg, drumstick and thigh, sauteed, ski...	0.0
7082	Industrial oil as ingredient in food	0.0

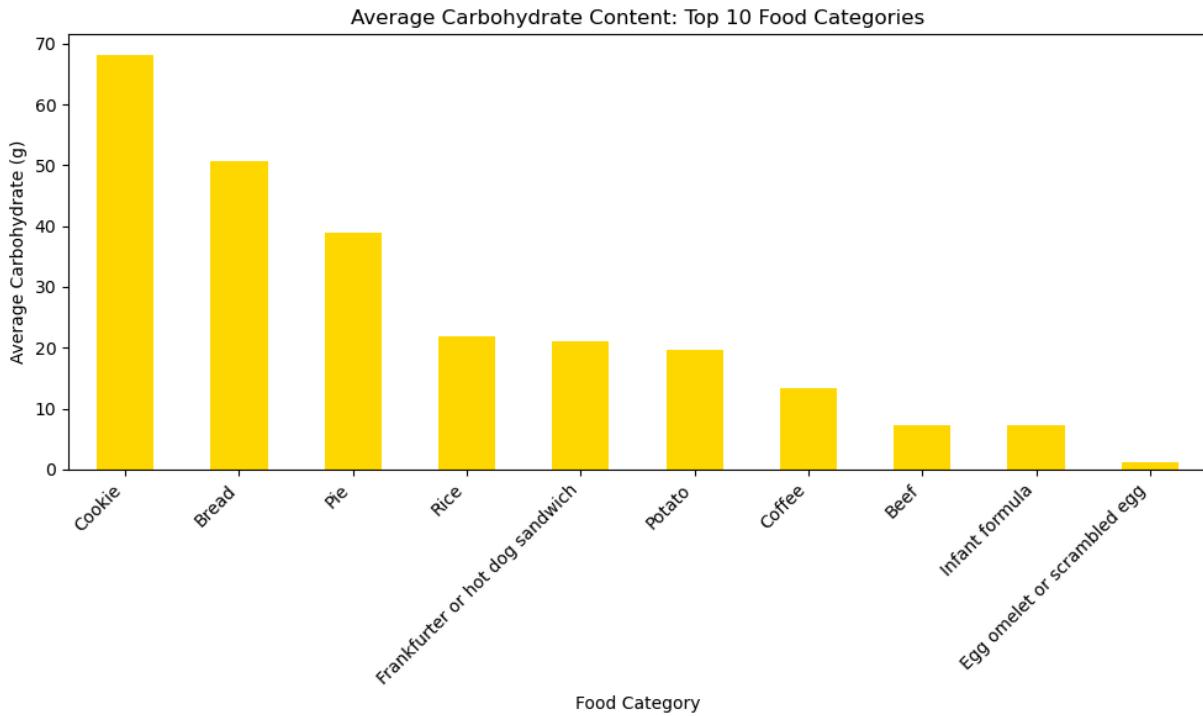
Mean Nutrient Content by Food Category

Bar charts below display the average protein, fat, and carbohydrate content for the ten most common food categories. Comparing these averages helps reveal broader nutrition trends and differences between food types.

```
In [9]: top_categories = data['Category'].value_counts().head(10).index
subset = data[data['Category'].isin(top_categories)]

avg_carb = subset.groupby('Category')['Data.Carbohydrate'].mean().sort_values(ascending=False)

plt.figure(figsize=(10,6))
avg_carb.plot(kind='bar', color='gold')
plt.title('Average Carbohydrate Content: Top 10 Food Categories')
plt.ylabel('Average Carbohydrate (g)')
plt.xlabel('Food Category')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



Interpretation: Nutritional Patterns by Food Category

By repeating this analysis for protein, fat, and carbohydrates, we can see how

different food types contribute to the macronutrient content of diets. High-protein foods cluster around animal products, while baked goods and processed foods tend to be higher in carbohydrates and/or fat. Visualizing by category helps identify trends that are useful for dietary planning and recommendation.

In [10]:

```
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

# Prepare features to cluster
top_categories = data['Category'].value_counts().head(10).index
subset = data[data['Category'].isin(top_categories)].copy() # <--- .copy() removes

features = subset[['Data.Protein', 'Data.Fat.Total Lipid', 'Data.Carbohydrate']].fi
scaler = StandardScaler()
X_scaled = scaler.fit_transform(features)

kmeans = KMeans(n_clusters=3, random_state=42)
subset['Cluster'] = kmeans.fit_predict(X_scaled)

subset[['Description', 'Category', 'Data.Protein', 'Data.Fat.Total Lipid', 'Data.Ca
```

Out[10]:

	Description	Category	Data.Protein	Data.Fat.Total Lipid	Data.Carbohydrate	Cluster
190	Infant formula, NFS	Infant formula	1.38	3.56	7.12	1
191	Infant formula, NS as to form (Similac Expert)	Infant formula	1.80	3.63	6.77	1
	...					
192	Infant formula, ready-to-feed (Similac Expert)	Infant formula	1.80	3.63	6.77	1
	...					
193	Infant formula, powder, made with water, NFS (...)	Infant formula	1.80	3.63	6.77	1
194	Infant formula, powder, made with tap water (S...)	Infant formula	1.80	3.63	6.77	1
195	Infant formula, powder, made with plain bottle...	Infant formula	1.80	3.63	6.77	1
196	Infant formula, powder, made with baby water (...)	Infant formula	1.80	3.63	6.77	1
197	Infant formula, NS as to form (Similac Advance)	Infant formula	1.37	3.62	6.87	1
198	Infant formula, ready-to-feed (Similac Advance)	Infant formula	1.37	3.62	6.87	1
199	Infant formula, liquid concentrate, made with ...	Infant formula	1.37	3.62	6.87	1

In this step, we prepare the food nutrition dataset for a clustering analysis. We select the ten most common food categories in

the data and focus on three key macronutrients: protein, fat, and carbohydrate. These features are standardized before using K-means clustering to assign each food to one of three nutrition-based clusters. This process helps uncover patterns in macronutrient profiles and allows us to group foods with similar nutritional characteristics.

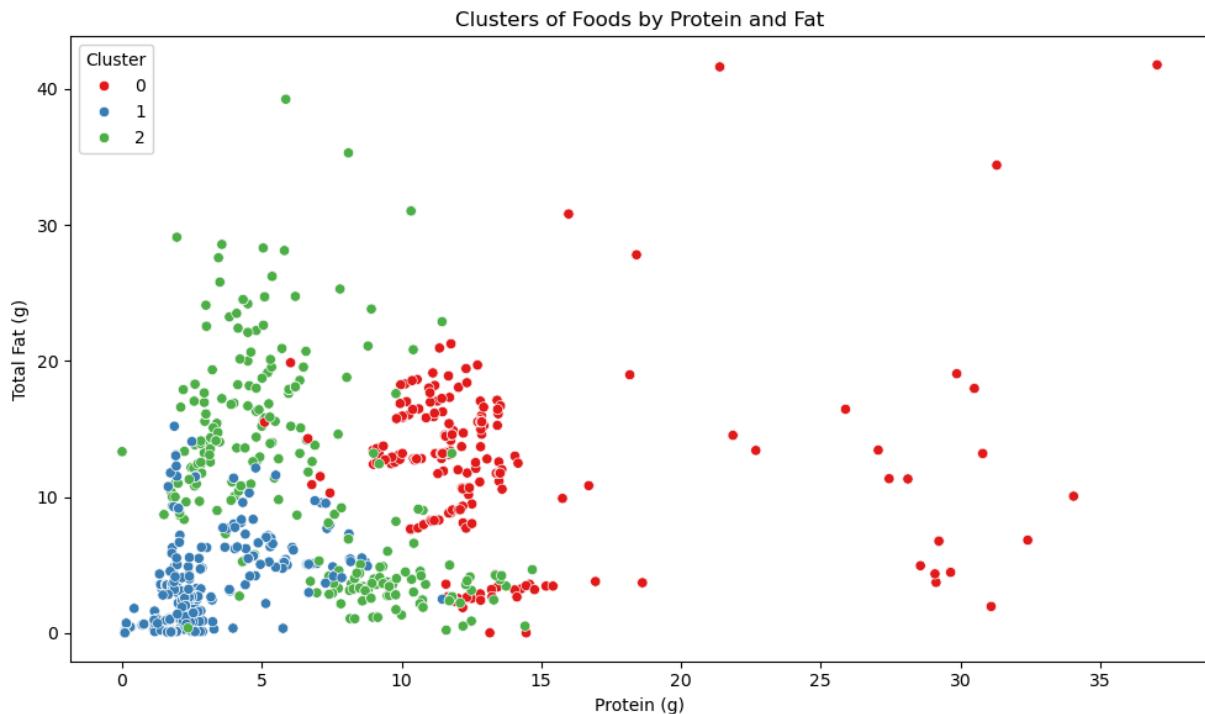
In [11]: `# Show foods from another cluster, e.g., cluster 0
subset[subset['Cluster'] == 0][['Description', 'Category', 'Data.Protein', 'Data.Fat.TotalLipid', 'Data.Carbohydrate']]`

Out[11]:

		Description	Category	Data.Protein	Data.Fat.TotalLipid	Data.Carbohydrate
649	Beef, NS as to cut, cooked, NS as to fat eaten	Beef		29.23	6.75	0.00
650	Beef, NS as to cut, cooked, lean and fat eaten	Beef		27.06	13.44	0.00
651	Beef, NS as to cut, cooked, lean only eaten	Beef		29.23	6.75	0.00
655	Beef, pickled	Beef		18.17	18.98	0.47
656	Beef, NS as to cut, fried, NS to fat eaten	Beef		29.23	6.75	0.00
675	Beef, oxtails, cooked	Beef		30.80	13.19	0.00
676	Beef, neck bones, cooked	Beef		29.87	19.06	0.00
677	Beef, shorribs, cooked, NS as to fat eaten	Beef		30.50	17.98	0.00
678	Beef, shorribs, cooked, lean and fat eaten	Beef		21.39	41.63	0.00
679	Beef, shorribs, cooked, lean only eaten	Beef		30.50	17.98	0.00

In [12]: `import seaborn as sns`

```
In [13]: plt.figure(figsize=(10,6))
sns.scatterplot(data=subset,
                 x='Data.Protein',
                 y='Data.Fat.Total Lipid',
                 hue='Cluster',
                 palette='Set1')
plt.title('Clusters of Foods by Protein and Fat')
plt.xlabel('Protein (g)')
plt.ylabel('Total Fat (g)')
plt.tight_layout()
plt.show()
```



The scatter plot above displays the results of K-means clustering on foods from the most common categories in the dataset, based on their protein and total fat content. Each point is a food, and its color represents the cluster assigned by the algorithm.

The visualization reveals distinct nutrient profiles one cluster primarily contains foods high in both protein and fat (often animal-based or processed), another

groups items with low protein and fat (such as some infant formulas or lighter foods), and the third cluster lies between these extremes. Clustering helps simplify and highlight complex dietary data, making it clear which foods are nutritionally similar. This approach is valuable for recommendations, menu planning, or understanding broad patterns in the food supply.

In []:

In []: